

### Goal 3: Turn my code (stock-rsi.py) into a containerized application, deploy it in Kubernetes (and prove it works!)

Turning a python program into a docker container is fairly straightforward. I used this guide as a starting point: <https://www.docker.com/blog/containerized-python-development-part-1/>

I built containers on my laptop and the resulting container works seamlessly on any other system that supports docker (including Kubernetes).

Here is a link to my build directory, including my Dockerfiles, and all dependent files: <https://github.com/olivelawn/cs199-indep-k8s/tree/main/container-stock-rsi>

A Dockerfile contains the instructions on how to build a container image. In my case, I needed to instruct docker to build an image that includes all of the modules (yfinance, pandas, lumpy, mysql.connector, etc) that my program relies on. For this there is a list of dependent modules in requirements.txt and a call to pip ("RUN pip install -r requirements.txt") to install them. Additionally, I needed to add into my image information on how to connect to the mysql db (config.py: connection string, credentials), add my file with the list of stock symbols (i.e. "symbols") and then finally instructions on how to execute the python program itself ("CMD ["python", "/app/stock-rsi.py", "symbols"]").

This command will build the docker container:

```
[parker@parker container-stock-rsi]$ docker build -f Dockerfile -t pjohn07/stock-rsi-k8s:latest .
[+] Building 3.0s (11/11) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 199B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/python:3.7
=> [auth] library/python:pull token for registry-1.docker.io
=> [1/5] FROM docker.io/library/python:3.7@sha256:354a4c2400dd1c1293b657bcf059742e45f7c6222dc908fc63ba4538740805e5
=> [internal] load build context
=> => transferring context: 154B
=> CACHED [2/5] RUN mkdir /app
=> CACHED [3/5] WORKDIR /app
=> CACHED [4/5] ADD . /app/
=> CACHED [5/5] RUN pip install -r requirements.txt
=> exporting to image
=> => exporting layers
=> => writing image sha256:a00ec65018021351d94858f52cff3dc2dc674036cfd1b29c9235c83c32a6e828
=> => naming to docker.io/pjohn07/stock-rsi-k8s:latest
```

(The '-t' option specifies a tag, which in my case is "pjohn07/stock-rsi-k8s" "pjohn07" is my username on dockerhub, which is where I store my images—and where Kubernetes looks by default to retrieve them)

Afterwards you can list the image and push it to dockerhub:

```
[parker@parker container-stock-rsi]$ docker image ls
REPOSITORY          TAG         IMAGE ID      CREATED       SIZE
pjohn07/stock-rsi-k8s  latest      a00ec6501802  40 hours ago  1.16GB
pjohn07/hello-python  latest      1aec00ffb622  2 days ago    887MB

[parker@parker container-stock-rsi]$ docker push pjohn07/stock-rsi-k8s
Using default tag: latest
The push refers to repository [docker.io/pjohn07/stock-rsi-k8s]
74332da7ee5d: Layer already exists
a8d3f57df352: Layer already exists
5f70bf18a086: Layer already exists
420960fa9cc9: Layer already exists
df58ce8478ec: Layer already exists
latest: digest: sha256:d14015656d597b3f0ac5b065e7aa3e716489daf3e797e3fd592a1b6123254b06 size: 3051
```

Here is stock-rsi.py on docker hub: <https://hub.docker.com/repository/docker/pjohn07/stock-rsi-k8s>

Next up is to deploy the container/pod into my Kubernetes cluster and insure it works.

At this point only mysql is running in the cluster.

```
parker@c1-master1:~$ kubectl get deployments.apps ; kubectl get svc
NAME      READY  UP-TO-DATE  AVAILABLE  AGE
mysql     1/1    1           1          18h
NAME      TYPE        CLUSTER-IP  EXTERNAL-IP  PORT(S)    AGE
mysqlsvc  ClusterIP   10.109.20.156 <none>       3306/TCP    18h
```

Currently there are 146 rows in our "stocks" table:

```
mysql> select COUNT(date) from stocks;
+-----+
| COUNT(date) |
+-----+
|          146 |
+-----+
1 row in set (0.03 sec)
```

Next I define a Kubernetes deployment for my stock-rsi container in yaml as seen here: <https://github.com/olivelawn/cs199-indep-k8s/blob/main/k8s-configs/stock-rsi-k8s/stock-rsi-deploy.yaml>

Notice I reference my docker hub image with "image: pjohn07/stock-rsi-k8s"

Let's deploy the container:

```
parker@c1-master1:~/stock-rsi-k8s$ kubectl apply -f stock-rsi-deploy.yaml
deployment.apps/stock-rsi created
```

Let's watch the pod startup and go from Pending, Creating, and Running status:

```
parker@c1-master1:~/stock-rsi-k8s$ while true; do kubectl get pods; sleep 10; done
NAME                                READY    STATUS    RESTARTS   AGE
mysql-68c777748-dvtjl              1/1      Running   1           19h
stock-rsi-6f778764b5-kj5jt         1/1      ContainerPending 0           8s
NAME                                READY    STATUS    RESTARTS   AGE
mysql-68c777748-dvtjl              1/1      Running   1           19h
stock-rsi-6f778764b5-kj5jt         1/1      ContainerCreating 1           18s
NAME                                READY    STATUS    RESTARTS   AGE
mysql-68c777748-dvtjl              1/1      Running   1           19h
stock-rsi-6f778764b5-kj5jt         1/1      Running   1           28s
```

As it turns out, inspecting stdout from inside the container is something of a challenge. It is possible, but you have to catch the container in just the right state (just after the container itself loads and just before python is kicked off) and attach a shell to the container with this command "kubectl exec -it stock-rsi-6f778764b5-kj5jt -- /bin/bash" (I tried several times without any luck). Instead, it is easier to have a look at the changes made to mysql "stocks" table.

```
mysql> select COUNT(date) from stocks;
+-----+
| COUNT(date) |
+-----+
|          147 |
+-----+
1 row in set (0.33 sec)
```

Notice our row count increased by 1 (from 146 to 147) Just to be sure, let's have a look at the data for today's date:

```
parker@c1-master1:~/mysql$ date
Fri May 21 15:47:45 UTC 2021
```

```
mysql> select * from stocks where date = '2021-05-21';
| date | AAGIY | AAPL | ABT | ADBE | ADDYY | AER | AES | AIG | AKAM | ALL | AMADY | AMAT | AMGN | AMZN |
ASML | ATVI | AVGO | AXAHY | AXP | BABA | BAC | BAESY | BASFY | BAX | BK | BKNG | BLK | BMY | BND | BNDX |
BUD | BURBY | C | CABGY | CAH | CARR | CAT | CHKP | CI | CMCSA | COP | COST | CRARY | CRM | CSCO | CSX | CVS |
CVX | DANoy | DBSDY | DD | DEO | DIS | DLTR | DNHBY | DOX | DPSGY | DUK | EADSY | EBKDY | EDU | EMB | EQR |
ESLOY | EW | F | FB | FIS | FMC | FTI | GLW | GOOG | GOOGL | GWW | HAL | HD | HON | HSBC | HTHY |
HUM | ICLR | IDEXY | IDXX | ILMN | INFO | INTC | IOO | ITOT | IWY | JBAXY | JETS | JNJ | JPM | LGLV | LGRDY |
LIN | LNSTY | LVMUY | LVS | LYG | MDLZ | MET | MO | MPC | MRK | MRO | MS | MSCI | MSFT | NEE | NGLOY | NKE |
NOC | NOW | NRDBY | NSRGY | NTN | NVDA | NVS | ORCL | OTIS | OXY | PEP | PFE | PG | PM | PNGAY | PRU |
PXD | PYPL | QCOM | QQQ | RCI | RDSMY | RHHBY | RTX | SAP | SCHW | SNN | SPLK | SQ | SSDOY | STNE | SU | T |
TCOM | TFC | TMO | TMSNY | TRYIY | TSM | TSN | TXN | UBER | UMICY | UNH | UNM | UNP | UPS | V | VDE |
VLO | VMW | VNQ | VNQI | VTI | VZ | WFC | WMB | ZTS | ABNB | MCHI | VB | VO | VTRS | VWO | ARKK |
CNRG | IWR | JKH | MSSMX | SCHD | IMCG | THNPFY | TDUP | ABBV | ACN | CHTR | DOW | EXC | GD | GS | IEMG |
KHC | LLY | LMT | MCD | NFLX | SCHF | SO | SPG | TFI | TGT | VEA | VOO | VTEB | VXF | WMT |
| 2021-05-21 | 55.73 | 46.06 | 49.5 | 50.76 | 71.19 | 46.91 | 42.05 | 62.11 | 68.99 | 70.48 | NULL | 51.92 | 55.43 | 44.75 |
51.5 | 56.85 | 50.48 | 48.03 | 62.41 | 37.68 | 63.03 | NULL | 45.86 | 46.23 | 60.71 | 47.01 | 61.99 | 70.77 | 50.09 | 39.91 |
68.58 | NULL | 65.49 | NULL | 48.3 | 56.82 | 55.04 | 53.63 | 61.8 | 47.44 | 55.11 | 59.81 | NULL | 54.1 | NULL | 44.08 |
81.19 | 47.82 | NULL | 58.04 | 66.2 | NULL | 32.09 | 38.39 | 55.18 | 55.28 | 65.44 | 57.98 | 51.31 | 80.44 | 27.72 | 53.91 |
57.51 | NULL | 63.61 | 66.61 | 58.38 | 53.28 | 54.04 | 55.03 | 46.62 | 57.71 | 55.99 | 61.67 | 55.43 | 44.62 | 50.18 | NULL |
NULL | 54.41 | 68.63 | NULL | 57.57 | 59.36 | NULL | 42.02 | 55.39 | 54.22 | 50.99 | NULL | 51.34 | 68.55 | 61.26 | 56.34 |
67.43 | 60.88 | NULL | 59.18 | NULL | NULL | 72.12 | 54.42 | 54.35 | NULL | 61.88 | 53.47 | 63.6 | 47.22 | 47.83 | 49.16 |
NULL | 47.66 | 66.35 | 43.17 | NULL | 68.75 | 62.12 | 56.56 | 54.36 | 61.83 | NULL | 48.87 | 60.94 | 65.12 | 61.31 | 60.79 |
33.61 | 65.29 | 47.19 | 50.53 | 47.27 | 49.98 | NULL | 52.81 | 61.33 | 65.05 | 55.21 | NULL | 64.74 | 44.94 | 36.7 | NULL |
NULL | 56.91 | 41.5 | 56.87 | 56.9 | 50.86 | 54.7 | 52.4 | 45.08 | 64.9 | 52.17 | 45.08 | NULL | 66.03 | 56.98 | 50.11 |
69.71 | 53.73 | 56.26 | 54.26 | NULL | 58.21 | 54.45 | 54.17 | 40.96 | 57.29 | 69.3 | 63.42 | 28.83 | 45.15 | 50.77 | 55.11 |
64 | 48.58 | 42.5 | 47.36 | 54.21 | NULL | 43.24 | 59.35 | 51.46 | 53.67 | 58.45 | 65.63 | 49.54 | 62 | 60.07 | 60.64 | 54.6 |
59.63 | 48.28 | 68.87 | 69.43 | 56.46 | 50.55 | 48.5 | 57.35 | 46.73 | 56.73 | 49.45 | 70.18 | 57.35 | 55.1 | 55.7 | 50.03 |
59.71 |
```

This proves that my container running my stock-rsi.py program is able to pull stock data from the internet, crunch some numbers, and insert that data into a separate container running mysql.